

Zhouxiang Fang

✉ zf28@rice.edu ☎ +1 (667) 421-9685 📍 Houston, TX

Research Interests: Large Language Models, LLM Safety and Alignment, NLP

EDUCATION

Rice University <i>Ph.D. in Computer Science</i>	Houston, TX <i>Aug 2025 – Present</i>
Johns Hopkins University <i>M.S.E. in Computer Science</i>	Baltimore, MD <i>Aug 2023 – Dec 2024</i>
Zhejiang University <i>B.E. in Artificial Intelligence (Turing Class, Chu Kochen Honors College)</i>	Hangzhou, China <i>Sep 2019 – June 2023</i>

RESEARCH EXPERIENCE

Chili Lab, Rice University <i>Ph.D. Student, supervised by Prof. Hanjie Chen</i>	Houston, TX <i>Aug 2025 – Present</i>
• GR-SAP (EMNLP'26): We proposed a unified framework to preserve LLM safety during fine-tuning without the typically inaccessible original alignment data, by replaying synthesized domain-specific substitutes.	
IALab, Johns Hopkins University <i>Research Assistant, supervised by Prof. Daniel Khashabi</i>	Baltimore, MD <i>June 2024 – May 2025</i>
• ICL CIPHERS (EMNLP'25): We proposed a class of task reformulations based on substitution ciphers to distinguish whether LLMs genuinely learn from in-context demonstrations or retrieve patterns seen in pre-training.	
Johns Hopkins University <i>Research Assistant, supervised by Prof. Mark Dredze</i>	Baltimore, MD <i>Sep 2023 – June 2024</i>
• Clinical QA Benchmarks (NAACL'25): We constructed two realistic and challenging clinical QA datasets with expert-written explanations to evaluate both LLMs' answers and their reasoning beyond licensing-exam questions.	
Zhejiang University <i>Research Assistant, supervised by Prof. Yang Yang</i>	Hangzhou, China <i>June 2022 – May 2023</i>
• Popular Headline Generation (AI Open'24): We proposed MEBART, a BART-based model that jointly captures evolving platform-wide trends and individual writing styles to generate popular social media headlines.	

PUBLICATIONS

1. **Zhouxiang Fang**, Jiawei Zhou, and Hanjie Chen. “GR-SAP: Generative Replay for Safety Alignment Preservation during Fine-Tuning.” In *EMNLP'26*, Oct, 2026.
2. **Zhouxiang Fang**, Aayush Mishra, Muhan Gao, Anqi Liu, and Daniel Khashabi. “ICL CIPHERS: Quantifying ‘Learning’ in In-Context Learning via Substitution Ciphers.” In *EMNLP'25*, pages 25912–25933, Nov, 2025.
3. Hanjie Chen*, **Zhouxiang Fang***, Yash Singla, and Mark Dredze. “Benchmarking Large Language Models on Answering and Explaining Challenging Medical Questions.” In *NAACL'25*, pages 3563–3599, Apr, 2025.
4. **Zhouxiang Fang**, Min Yu, Zhendong Fu, Boning Zhang, Xuanwen Huang, Xiaoqi Tang, and Yang Yang. “How to Generate Popular Post Headlines on Social Media?” In *AI Open'24 (Vol 5)*, pages 1–9, 2024.

*Equal contribution.

SKILLS

Programming: C/C++, Python, HTML, CSS, JavaScript, Julia